

# Sentiment Analysis on Airline Tweets Dataset

The background is a dark blue gradient. On the left, there are several interlocking gears. One large gear in the center contains the text 'NLP Natural Language Processing'. Other gears contain icons: a robot head with a speech bubble, two speech bubbles, and a head profile with gears inside. A hand from the right side of the frame points towards the center. In the top left, there are faint labels 'CRM' and 'Quality'. At the bottom, there are faint icons of a group of people and a world map.

Steve Song

Aug 12 ,2020

# Problem Statement

- Twitter as a part of the marketing strategy for business owners, is a gold mine of customer insights and opportunities to build the brand, drive sales and win fans.
- The dataset we acquired contains tweets on US Airline of February 2015 classified in positive, negative and neutral tweets.
- By applying the sentiment analysis machine learning techniques, we build the predictive models, which enable business owners to predict whether the customers' feedback is positive, negative or neutral.

# Data pre-processing

- Tweet preprocessor:
  - a preprocessing library for tweet data written in Python
  - supports cleaning, tokenizing and parsing: URLs, Hashtags, Mentions, Reserved words (RT,FAV),Emojis and Smileys.
  - For example, the original content of the first tweet is **@VirginAmerica What @dhepburn said.**
  - it was converted into “What said”.

# Data pre-processing

- Expanding contractions:
  - Converting each contraction to its expanded, original form often helps with text standardization
  - For example, is “@VirginAmerica I didn't today... Must mean I need to take another trip!. The step converted “didn’t “ into “did not”.

# Data pre-processing

- Removing special characters:
  - Special characters and symbols which are usually non alphanumeric characters often add to the extra noise in unstructured text.
  - In this project, we use the regular expression `'[^a-zA-z0-9\s]'`.

# Data pre-processing

## Stemming and lemmatization


- Word stems are usually the base form of possible words that can be created by attaching *affixes* like *prefixes* and *suffixes* to the stem to create new words.
- For example, the words “said” have been converted into “say”.

# Data pre-processing

Removing stopwords:

- Words which have little or no significance especially when constructing meaningful features from text are known as stopwords or stop words.
- We use package NLTK to create a stopwords list, which contains 17 words.


# Exploratory Data Analysis (EDA)



this dataset doesn't have an issue of 'imbalance'.

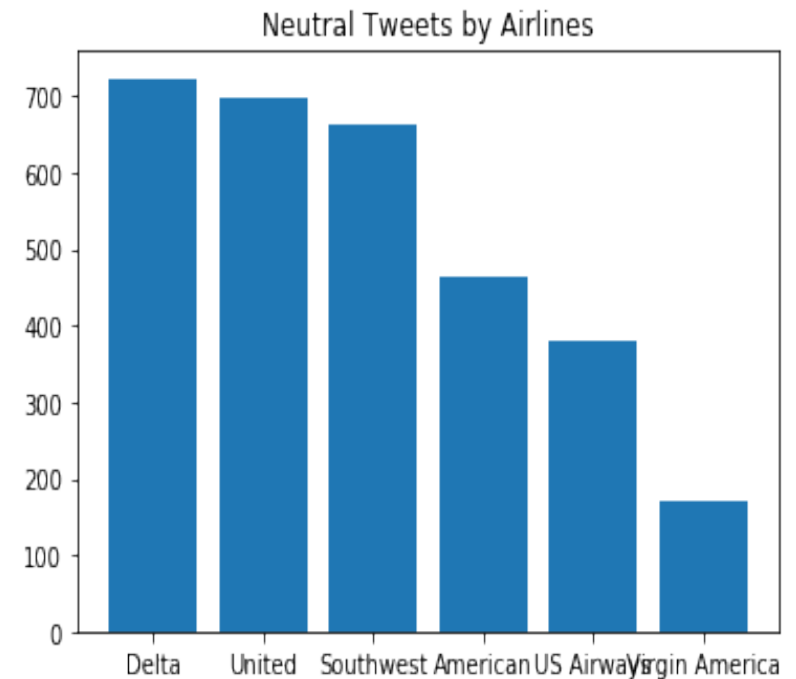
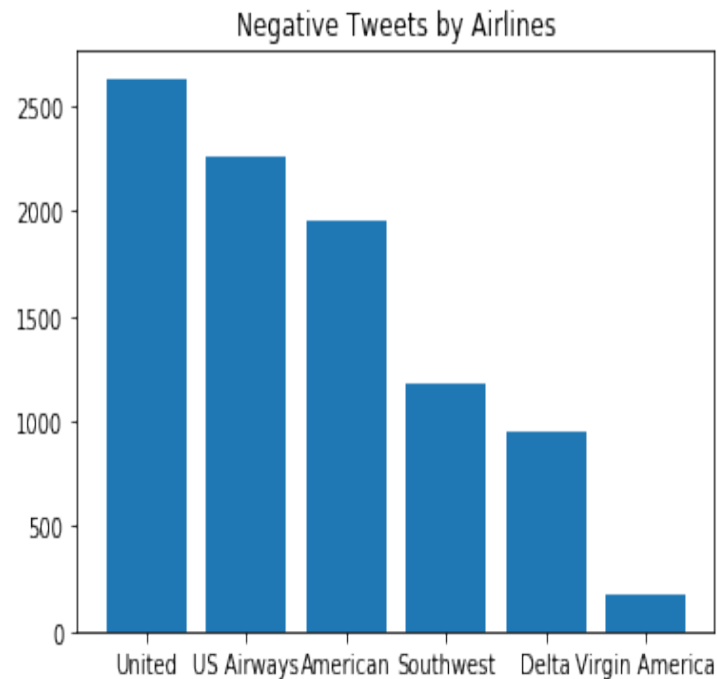
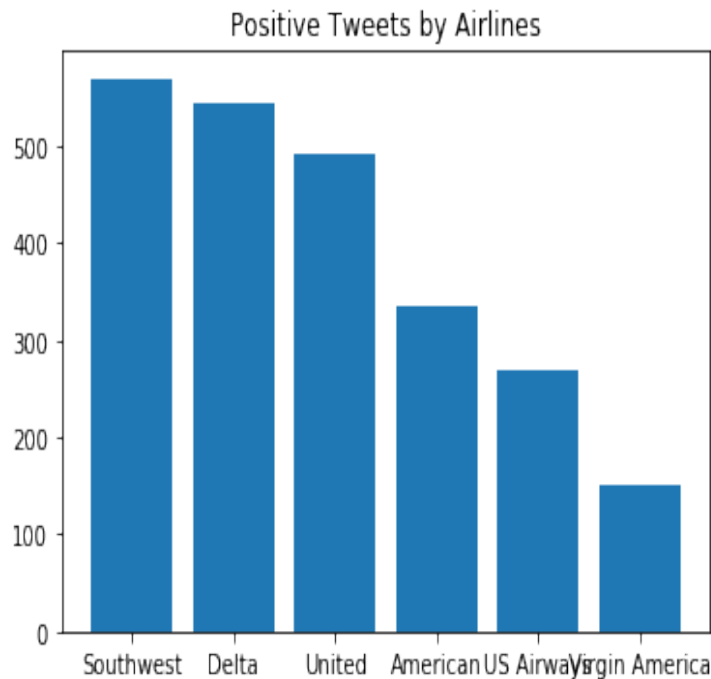


# Exploratory Data Analysis (EDA)



top 3 reasons are Customer Service, Late Flight and Can't tell.

# Exploratory Data Analysis (EDA)



the Southwest had the best performance in terms of positive tweets;  
United had the worst performance measured by negative tweets.

# Exploratory Data Analysis (EDA)



The most frequently occurred words are '@', '.', 'to', 'I', 'the', '!', '?', 'a', ',', 'for'.

# Text data encoding

## Word Counts with CountVectorizer

|   | 00 | 000 | 000419 | 00a | 00am | 00p | 00pm | 0200 | 03 | 04 | ... | zagg | zambia | zcc82u | zero | zig | zip | zipper | zone | zoom | zurich |
|---|----|-----|--------|-----|------|-----|------|------|----|----|-----|------|--------|--------|------|-----|-----|--------|------|------|--------|
| 0 | 0  | 0   | 0      | 0   | 0    | 0   | 0    | 0    | 0  | 0  | ... | 0    | 0      | 0      | 0    | 0   | 0   | 0      | 0    | 0    | 0      |
| 1 | 0  | 0   | 0      | 0   | 0    | 0   | 0    | 0    | 0  | 0  | ... | 0    | 0      | 0      | 0    | 0   | 0   | 0      | 0    | 0    | 0      |
| 2 | 0  | 0   | 0      | 0   | 0    | 0   | 0    | 0    | 0  | 0  | ... | 0    | 0      | 0      | 0    | 0   | 0   | 0      | 0    | 0    | 0      |
| 3 | 0  | 0   | 0      | 0   | 0    | 0   | 0    | 0    | 0  | 0  | ... | 0    | 0      | 0      | 0    | 0   | 0   | 0      | 0    | 0    | 0      |
| 4 | 0  | 0   | 0      | 0   | 0    | 0   | 0    | 0    | 0  | 0  | ... | 0    | 0      | 0      | 0    | 0   | 0   | 0      | 0    | 0    | 0      |

5 rows x 6875 columns

we created a data frame with 6875 columns.

# Text data encoding

## Word Frequencies with TfidfVectorizer

|   | 00  | 000 | 000419 | 00a | 00am | 00p | 00pm | 0200 | 03  | 04  | ... | zagg | zambia | zcc82u | zero | zig | zip | zipper | zone | zoom | zurich |
|---|-----|-----|--------|-----|------|-----|------|------|-----|-----|-----|------|--------|--------|------|-----|-----|--------|------|------|--------|
| 0 | 0.0 | 0.0 | 0.0    | 0.0 | 0.0  | 0.0 | 0.0  | 0.0  | 0.0 | 0.0 | ... | 0.0  | 0.0    | 0.0    | 0.0  | 0.0 | 0.0 | 0.0    | 0.0  | 0.0  | 0.0    |
| 1 | 0.0 | 0.0 | 0.0    | 0.0 | 0.0  | 0.0 | 0.0  | 0.0  | 0.0 | 0.0 | ... | 0.0  | 0.0    | 0.0    | 0.0  | 0.0 | 0.0 | 0.0    | 0.0  | 0.0  | 0.0    |
| 2 | 0.0 | 0.0 | 0.0    | 0.0 | 0.0  | 0.0 | 0.0  | 0.0  | 0.0 | 0.0 | ... | 0.0  | 0.0    | 0.0    | 0.0  | 0.0 | 0.0 | 0.0    | 0.0  | 0.0  | 0.0    |
| 3 | 0.0 | 0.0 | 0.0    | 0.0 | 0.0  | 0.0 | 0.0  | 0.0  | 0.0 | 0.0 | ... | 0.0  | 0.0    | 0.0    | 0.0  | 0.0 | 0.0 | 0.0    | 0.0  | 0.0  | 0.0    |
| 4 | 0.0 | 0.0 | 0.0    | 0.0 | 0.0  | 0.0 | 0.0  | 0.0  | 0.0 | 0.0 | ... | 0.0  | 0.0    | 0.0    | 0.0  | 0.0 | 0.0 | 0.0    | 0.0  | 0.0  | 0.0    |

5 rows x 6875 columns

# Traditional Machine Learning Model Building

## -- CountVectorization

**Naïve Bayes multinomial  
classification**

**Accuracy: 0.7491**

**Support Vector  
classification**

**Accuracy: 0.7692**

## -- TF-IDF

**Naïve Bayes multinomial  
classification**

**Accuracy: 0.6841**

**Support Vector  
classification**

**Accuracy: 0.7651**

# Using word embeddings

Word2vec is one of the most popular technique to learn word embeddings using a two-layer neural network.

Gensim library will enable us to develop word embeddings by training our own word2vec models on a custom corpus either with CBOW or skip-grams algorithms.

```
num_features = 100
min_word_count = 5
num_workers = multiprocessing.cpu_count()
context_size = 5
seed = 1
```

# Using word embeddings

**Logistic Regression**

**Accuracy: 0.7485**

**Support Vector  
classification**

**Accuracy: 0.7593**

- Gensim also enables us to use pre-trained word2vec model
- we will import pre-trained Google News model with 300 dimensions for word vector.

| SVC    | decision tree | logistic<br>regression |
|--------|---------------|------------------------|
| 0.7823 | 0.6226        | 0.7659                 |



# Deep learning on Keras

- Learning word embeddings with the embedding layer
  - for example , the sequences of word index for the first five documents in the corpus.

```
[[35],  
 [389, 186, 969, 99, 4428],  
 [2, 45, 561, 229, 23, 34, 84, 102],  
 [73, 2703, 2307, 3267, 753, 1569, 860, 24, 349, 2038],  
 [73, 296, 39, 177]]
```

---

# Deep learning on Keras

- the total number of unique tokens is 8,535

| Layer (type)              | Output Shape   | Param # |
|---------------------------|----------------|---------|
| embedding_1 (Embedding)   | (None, 22, 64) | 546304  |
| flatten_1 (Flatten)       | (None, 1408)   | 0       |
| dense_1 (Dense)           | (None, 3)      | 4227    |
| Total params: 550,531     |                |         |
| Trainable params: 550,531 |                |         |
| Non-trainable params: 0   |                |         |
| None                      |                |         |

# Deep learning on Keras

Train on 11712 samples, validate on 2928 samples

Epoch 1/10

11712/11712 [=====] - 7s 632us/step - loss: 0.8143 - acc: 0.6535 - val\_loss: 0.5426 - val\_acc: 0.7971

Epoch 2/10

11712/11712 [=====] - 4s 378us/step - loss: 0.5394 - acc: 0.7903 - val\_loss: 0.4743 - val\_acc: 0.8224

Epoch 3/10

11712/11712 [=====] - 5s 427us/step - loss: 0.3999 - acc: 0.8558 - val\_loss: 0.4590 - val\_acc: 0.8221

Epoch 4/10

11712/11712 [=====] - 5s 417us/step - loss: 0.3059 - acc: 0.8968 - val\_loss: 0.4729 - val\_acc: 0.8159

Epoch 5/10

11712/11712 [=====] - 5s 407us/step - loss: 0.2355 - acc: 0.9254 - val\_loss: 0.4894 - val\_acc: 0.8101

Epoch 6/10

11712/11712 [=====] - 5s 407us/step - loss: 0.1817 - acc: 0.9476 - val\_loss: 0.5189 - val\_acc: 0.8105

Epoch 7/10

11712/11712 [=====] - 5s 465us/step - loss: 0.1441 - acc: 0.9605 - val\_loss: 0.5489 - val\_acc: 0.8070

Epoch 8/10

11712/11712 [=====] - 7s 568us/step - loss: 0.1174 - acc: 0.9684 - val\_loss: 0.5743 - val\_acc: 0.8060

Epoch 9/10

11712/11712 [=====] - 6s 540us/step - loss: 0.0973 - acc: 0.9743 - val\_loss: 0.6027 - val\_acc: 0.8012

Epoch 10/10

11712/11712 [=====] - 6s 506us/step - loss: 0.0828 - acc: 0.9805 - val\_loss: 0.6303 - val\_acc: 0.8005

# Deep learning on Keras

- Using pre-trained word embedding, importing 100-dimensional “Glove” database, we created the dictionary of embedding index.
- 

| Layer (type)              | Output Shape    | Param # |
|---------------------------|-----------------|---------|
| =====                     |                 |         |
| embedding_3 (Embedding)   | (None, 22, 100) | 865200  |
| =====                     |                 |         |
| flatten_3 (Flatten)       | (None, 2200)    | 0       |
| =====                     |                 |         |
| dense_4 (Dense)           | (None, 32)      | 70432   |
| =====                     |                 |         |
| dense_5 (Dense)           | (None, 3)       | 99      |
| =====                     |                 |         |
| Total params: 935,731     |                 |         |
| Trainable params: 935,731 |                 |         |
| Non-trainable params: 0   |                 |         |
| =====                     |                 |         |

# Deep learning on Keras

Train on 11712 samples, validate on 2928 samples

```
Epoch 1/10
11712/11712 [=====] - 6s 495us/step - loss: 0.7063 - acc: 0.7130 - val_loss: 0.5474 - val_ac
c: 0.7889
Epoch 2/10
11712/11712 [=====] - 5s 423us/step - loss: 0.5741 - acc: 0.7693 - val_loss: 0.5228 - val_ac
c: 0.8057
Epoch 3/10
11712/11712 [=====] - 5s 426us/step - loss: 0.5090 - acc: 0.7989 - val_loss: 0.5450 - val_ac
c: 0.7951
Epoch 4/10
11712/11712 [=====] - 5s 424us/step - loss: 0.4494 - acc: 0.8250 - val_loss: 0.5787 - val_ac
c: 0.7862
Epoch 5/10
11712/11712 [=====] - 5s 433us/step - loss: 0.3941 - acc: 0.8502 - val_loss: 0.5863 - val_ac
c: 0.7913
Epoch 6/10
11712/11712 [=====] - 5s 423us/step - loss: 0.3422 - acc: 0.8727 - val_loss: 0.6409 - val_ac
c: 0.7746
Epoch 7/10
11712/11712 [=====] - 5s 431us/step - loss: 0.2927 - acc: 0.8965 - val_loss: 0.6780 - val_ac
c: 0.7671
Epoch 8/10
11712/11712 [=====] - 5s 429us/step - loss: 0.2515 - acc: 0.9107 - val_loss: 0.7410 - val_ac
c: 0.7408
Epoch 9/10
11712/11712 [=====] - 5s 424us/step - loss: 0.2127 - acc: 0.9296 - val_loss: 0.7528 - val_ac
c: 0.7835
Epoch 10/10
11712/11712 [=====] - 5s 422us/step - loss: 0.1820 - acc: 0.9409 - val_loss: 0.8168 - val_ac
c: 0.7698
```

# Deep learning on Keras

- Keras also provides two different layers, SimpleRNN & LSTM. We can use them together with embedding layer to construct neural network model.

```
Train on 11712 samples, validate on 2928 samples
Epoch 1/10
11712/11712 [=====] - 6s 518us/step - loss: 0.7643 - acc: 0.6867 - val_loss: 0.5623 - val_ac
c: 0.7964
Epoch 2/10
11712/11712 [=====] - 5s 394us/step - loss: 0.5662 - acc: 0.7836 - val_loss: 0.5132 - val_ac
c: 0.8012
Epoch 3/10
11712/11712 [=====] - 5s 389us/step - loss: 0.4733 - acc: 0.8227 - val_loss: 0.5505 - val_ac
c: 0.7862
Epoch 4/10
11712/11712 [=====] - 5s 400us/step - loss: 0.4126 - acc: 0.8487 - val_loss: 0.6630 - val_ac
c: 0.7377
Epoch 5/10
11712/11712 [=====] - 5s 390us/step - loss: 0.3578 - acc: 0.8725 - val_loss: 0.5906 - val_ac
c: 0.7848
Epoch 6/10
11712/11712 [=====] - 5s 393us/step - loss: 0.3133 - acc: 0.8881 - val_loss: 0.5705 - val_ac
c: 0.7900
Epoch 7/10
11712/11712 [=====] - 5s 399us/step - loss: 0.2772 - acc: 0.9016 - val_loss: 0.5892 - val_ac
c: 0.7937
Epoch 8/10
11712/11712 [=====] - 5s 407us/step - loss: 0.2402 - acc: 0.9152 - val_loss: 0.7106 - val_ac
c: 0.7490
Epoch 9/10
11712/11712 [=====] - 5s 422us/step - loss: 0.2084 - acc: 0.9263 - val_loss: 0.6624 - val_ac
c: 0.7702
Epoch 10/10
11712/11712 [=====] - 5s 424us/step - loss: 0.1825 - acc: 0.9365 - val_loss: 0.6946 - val_ac
c: 0.7811
```

# Deep learning on Keras

Train on 11712 samples, validate on 2928 samples

Epoch 1/10

11712/11712 [=====] - 15s 1ms/step - loss: 0.8327 - acc: 0.6361 - val\_loss: 0.5610 - val\_acc: 0.7688

Epoch 2/10

11712/11712 [=====] - 18s 2ms/step - loss: 0.6253 - acc: 0.7229 - val\_loss: 0.5160 - val\_acc: 0.7971

Epoch 3/10

11712/11712 [=====] - 20s 2ms/step - loss: 0.5508 - acc: 0.7717 - val\_loss: 0.5224 - val\_acc: 0.7900

Epoch 4/10

11712/11712 [=====] - 31s 3ms/step - loss: 0.5086 - acc: 0.7978 - val\_loss: 0.5100 - val\_acc: 0.7917

Epoch 5/10

11712/11712 [=====] - 26s 2ms/step - loss: 0.4782 - acc: 0.8147 - val\_loss: 0.5026 - val\_acc: 0.8033

Epoch 6/10

11712/11712 [=====] - 23s 2ms/step - loss: 0.4554 - acc: 0.8262 - val\_loss: 0.5226 - val\_acc: 0.8053

Epoch 7/10

11712/11712 [=====] - 24s 2ms/step - loss: 0.4338 - acc: 0.8336 - val\_loss: 0.5225 - val\_acc: 0.7848

Epoch 8/10

11712/11712 [=====] - 24s 2ms/step - loss: 0.4145 - acc: 0.8459 - val\_loss: 0.5235 - val\_acc: 0.8023

Epoch 9/10

11712/11712 [=====] - 28s 2ms/step - loss: 0.3908 - acc: 0.8595 - val\_loss: 0.5188 - val\_acc: 0.8016

Epoch 10/10

11712/11712 [=====] - 27s 2ms/step - loss: 0.3643 - acc: 0.8711 - val\_loss: 0.5492 - val\_acc: 0.7934

# Business Implication

- Out of all these models, we found the Keras-based model using Embedding layer outperforms other counterparts by the accuracy score of 0.8224.
- If this model is put into practice, the business owner will have the much high ability of predicting the sentiment for incoming messages. This can definitely help enhance the operational efficiency and improve the customers' satisfaction.